


Fresh width128: training is nearly flat; evaluation saturation is unmeasured


1° resolution · residual MP2 · Mamba1 di16 / BCG1 · batch 1 · seed 22 · 10,000 updates

Late training improves slightly. The final model beats early SWA, but two artifacts cannot establish an eval plateau.


A Recorded training loss




B Compare repeated data positions



C Actual learning-rate schedule



D Only two completed evaluation artifacts



Training loss averages all 24 BPTT steps on evolving closed-loop trajectories; it is distinct from 10-day evaluation loss.

Panel B controls data-position coverage; the last partial pass must not be compared directly with a full pass.

No evaluation at individual 2k-9k checkpoints was found. Better final performance than SWA does not establish continuing late eval gains.