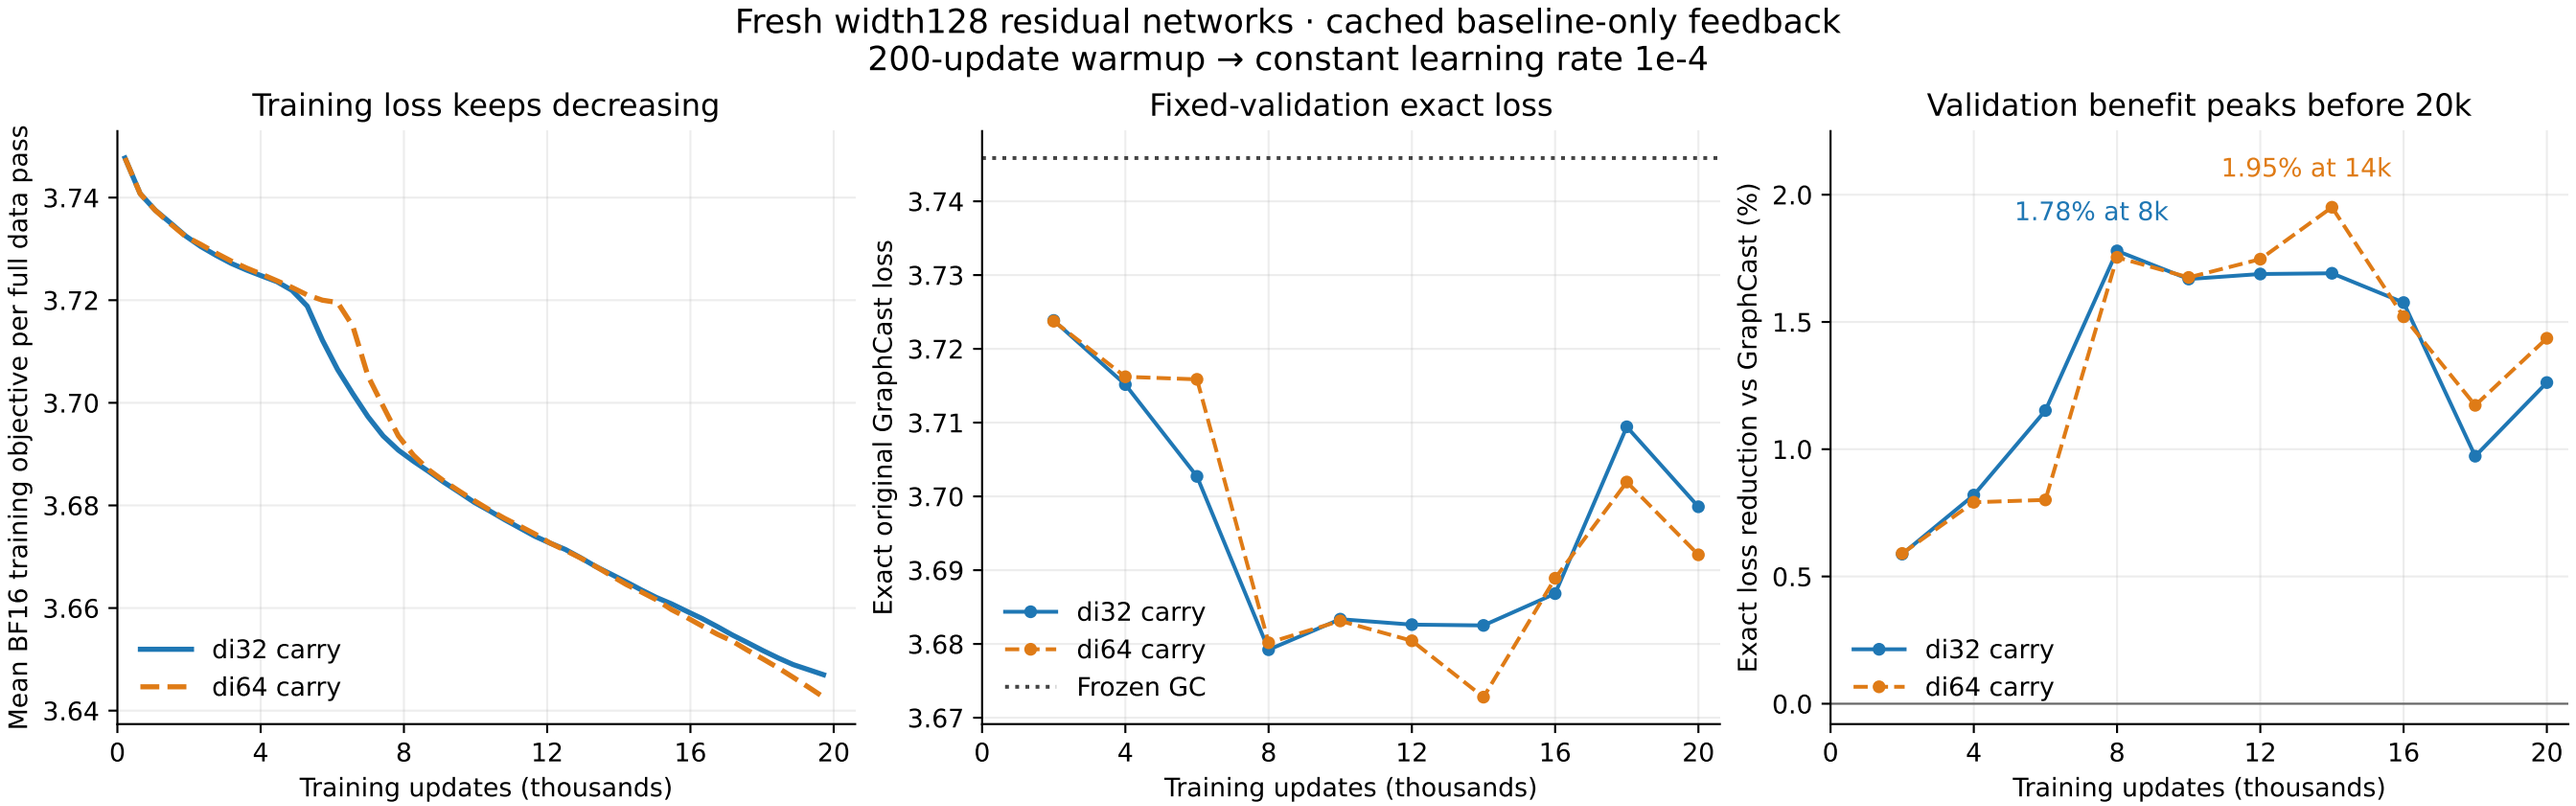
Fresh width128 residual networks · cached baseline-only feedback
200-update warmup → constant learning rate 1e-4
Training loss keeps decreasing
Mean BF16 training objective per full data pass
3.74
3.72
3.70
3.68
3.66
3.64
di32 carry
di64 carry
0
4
8
12
16
20
Training updates (thousands)
Fixed-validation exact loss
Exact original GraphCast loss
3.74
3.73
3.72
3.71
3.70
3.69
3.68
3.67
di32 carry
di64 carry
Frozen GC
0
4
8
12
16
20
Training updates (thousands)
Validation benefit peaks before 20k
Exact loss reduction vs GraphCast (%)
2.0
1.5
1.0
0.5
0.0
1.78% at 8k
1.95% at 14k
di32 carry
di64 carry
0
4
8
12
16
20
Training updates (thousands)
Training: 424 updates per complete pass; last partial pass omitted. Validation: same 60 chunks, all 24 positions including prefix.