


Res1: internal-state reset · approximate GraphCast loss


— Mamba: carry train / carry eval - - Mamba: carry train / reset eval - · Mamba: reset train / reset eval

Approx. GraphCast loss reduction (%)


di16 / BCG1 · SWA 2k-8k




di16 / BCG1 · SWA 4k-10k




di16 / BCG1 · SWA 6k-12k




di16 / BCG4 · SWA 2k-8k



di16 / BCG4 · SWA 4k-10k



di16 / BCG4 · SWA 6k-12k



Forecast lead time (days)

Forecast lead time (days)

Forecast lead time (days)

Relative to GraphCast-small baseline · 32 matched cold samples · 40 six-hour leads
Reset training = reset_every_anchor; reset evaluation = every step. Separate training runs.
Each curve uses its saved baseline (maximum difference 0.56%). Channel-MSE reconstruction; exact pole weighting unavailable.