







Does the trained model improve accuracy? All nonzero spatial modes

Positive = better than original GC · Negative = worse · Zero = original GC





— di16 · faster Mamba LR — di32 · joint LR — di16 · optimizer study

1° · Eight 2023 starts · 30–60° both hemispheres · Each model uses its own GC baseline.

Diagnostics: baseline repeatability unresolved; no uncertainty intervals. Spectral-error reduction is not exact GraphCast loss reduction.

Spectral-error reduction over the full 10-day rollout

Positive = better than original GC · Negative = worse · Zero = original GC



di16 · faster Mamba LR di32 · joint LR di16 · optimizer study

1° · Eight 2023 starts · 30–60° both hemispheres · Each model uses its own GC baseline.

Diagnostics: baseline repeatability unresolved; no uncertainty intervals. Spectral-error reduction is not exact GraphCast loss reduction.